

1. (a) Compile and run the program using the gcc options `-fstack-protector` and `-fstack-protector-all`. These options are used to detect buffer overflows when the program is being run. Explain why these options may not work with the program `bufOverflow.c` and how they can be enabled. (Hint: explore the use of the GCC option `--param=ssp-buffer-size=`)

Answer 1)

The `-fstack-protector` flag, and `-fstack-protector-all` flag, protect functions against stack smashing attacks and buffer overflows.

`-fstack-protector` adds a guard variable to functions with vulnerable objects. This includes functions that have buffers larger than 8 bytes. If a guard check fails, an error message is printed and the program exits. `-fstack-protector-all` does the same task as `-fstack-protector` but it protects all functions of the program.

While compiling `bufOverflow.c` the command to use the flags are

```
$ gcc -fstack-protector -g -o bof bufOverflow.c
```

```
or $ gcc -fstack-protector-all -g -o bof bufOverflow.c
```

For the method `bufOverflow.c` the char `buf[4]` is 4 bytes.

Hence the use of the above flags in compiling does not prove effective.

```
root@seed-desktop:/home/seed# gcc -fstack-protector -g -o bof bufOverflow.c
```

```
root@seed-desktop:/home/seed# ./bof < input
```

Enter the data

Should not reach here

```
root@seed-desktop:/home/seed# gcc -fstack-protector-all -g -o bof bufOverflow.c
```

```
root@seed-desktop:/home/seed# ./bof < input
```

Enter the data

Should not reach here

```
root@seed-desktop:/home/seed# gcc -fstack-protector --param ssp-buffer-size=4 -g -o bof bufOverflow.c
```

```
root@seed-desktop:/home/seed# ./bof < input
```

Enter the data

```
*** stack smashing detected ***: ./bof terminated
```

```
===== Backtrace: =====
```

```
/lib/tls/i686/cmov/libc.so.6(__fortify_fail+0x48)[0xb7f6cda8]
```

```
/lib/tls/i686/cmov/libc.so.6(__fortify_fail+0x0)[0xb7f6cd60]
```

```
./bof[0x8048553]
```

```
./bof[0x80485a1]
```

```
===== Memory map: =====
```

```
08048000-08049000 r-xp 00000000 08:01 8445 /home/seed/bof
```

```
08049000-0804a000 r--p 00000000 08:01 8445 /home/seed/bof
```

```
0804a000-0804b000 rw-p 00001000 08:01 8445 /home/seed/bof
```

```
0804b000-0806c000 rw-p 0804b000 00:00 0 [heap]
```

```
b7e52000-b7e5f000 r-xp 00000000 08:01 278049 /lib/libgcc_s.so.1
```

```
b7e5f000-b7e60000 r--p 0000c000 08:01 278049 /lib/libgcc_s.so.1
```

```
b7e60000-b7e61000 rw-p 0000d000 08:01 278049 /lib/libgcc_s.so.1
```

```
b7e6e000-b7e6f000 rw-p b7e6e000 00:00 0
```

```
b7e6f000-b7fcb000 r-xp 00000000 08:01 295506 /lib/tls/i686/cmov/libc-2.9.so
```

```
b7fcb000-b7fcc000 ---p 0015c000 08:01 295506 /lib/tls/i686/cmov/libc-2.9.so
```

```
b7fcc000-b7fce000 r--p 0015c000 08:01 295506 /lib/tls/i686/cmov/libc-2.9.so
```

```
b7fce000-b7fcf000 rw-p 0015e000 08:01 295506 /lib/tls/i686/cmov/libc-2.9.so
```

```

b7fcf000-b7fd2000 rw-p b7fcf000 00:00 0
b7fdd000-b7fe1000 rw-p b7fdd000 00:00 0
b7fe1000-b7fe2000 r-xp b7fe1000 00:00 0      [vdso]
b7fe2000-b7ffe000 r-xp 00000000 08:01 278007   /lib/ld-2.9.so
b7ffe000-b7fff000 r--p 0001b000 08:01 278007   /lib/ld-2.9.so
b7fff000-b8000000 rw-p 0001c000 08:01 278007   /lib/ld-2.9.so
bffe000-c0000000 rw-p bffe0000 00:00 0      [stack]
Aborted

```

For buffer overflow protection with `-fstack-protector` the `ssp-buffer-size` parameter can be set to the minimum size of buffer that requires protection. In this case the value will be 4.

```

root@seed-desktop: /home/seed
File Edit View Terminal Help
root@seed-desktop:/home/seed# ./bof < input
Enter the data
Should not reach here
root@seed-desktop:/home/seed# gcc -fstack-protector --param ssp-buffer-size=4 -g -o bof bufOverflow.c
root@seed-desktop:/home/seed# ./bof < input
Enter the data
*** stack smashing detected ***: ./bof terminated
===== Backtrace: =====
/lib/tls/i686/cmov/libc.so.6(__fortify_fail+0x48)[0xb7f6cda8]
/lib/tls/i686/cmov/libc.so.6(__fortify_fail+0x0)[0xb7f6cd60]
./bof[0x8048553]
./bof[0x80485a1]
===== Memory map: =====
08048000-08049000 r-xp 00000000 08:01 8445 /home/seed/bof
08049000-0804a000 r--p 00000000 08:01 8445 /home/seed/bof
0804a000-0804b000 rw-p 00001000 08:01 8445 /home/seed/bof
0804b000-0806c000 rw-p 0804b000 00:00 0 [heap]
b7e52000-b7e5f000 r-xp 00000000 08:01 278049 /lib/libgcc_s.so.1
b7e5f000-b7e60000 r--p 0000c000 08:01 278049 /lib/libgcc_s.so.1
b7e60000-b7e61000 rw-p 0000d000 08:01 278049 /lib/libgcc_s.so.1
b7e6e000-b7e6f000 rw-p b7e6e000 00:00 0
b7e6f000-b7fcb000 r-xp 00000000 08:01 295506 /lib/tls/i686/cmov/libc-2.9.so
b7fcb000-b7fcc000 ---p 0015c000 08:01 295506 /lib/tls/i686/cmov/libc-2.9.so
b7fcc000-b7fce000 r--p 0015c000 08:01 295506 /lib/tls/i686/cmov/libc-2.9.so
b7fce000-b7fcf000 rw-p 0015e000 08:01 295506 /lib/tls/i686/cmov/libc-2.9.so
b7fcf000-b7fd2000 rw-p b7fcf000 00:00 0
b7fdd000-b7fe1000 rw-p b7fdd000 00:00 0
b7fe1000-b7fe2000 r-xp b7fe1000 00:00 0 [vdso]
b7fe2000-b7ffe000 r-xp 00000000 08:01 278007 /lib/ld-2.9.so
b7ffe000-b7fff000 r--p 0001b000 08:01 278007 /lib/ld-2.9.so
b7fff000-b8000000 rw-p 0001c000 08:01 278007 /lib/ld-2.9.so
bffe000-c0000000 rw-p bffe0000 00:00 0 [stack]
Aborted
root@seed-desktop:/home/seed#

```

(b) Draw a diagram of the stack showing how it looks

- immediately before the `strcpy()` function is executed. You must show where the argument `str`, saved base pointer of `main`, return address of `main()` and local variable `buf` are placed on the stack.
- immediately after the `strcpy()` function completes. Show the contents of the buffer and how the input to the program is stored on the stack.

Answer part a)

The `malinput.c` file dumps text data, to be used with malicious intent in the `bufOverflow` program, in the input file.

```
seed@seed-desktop:~$ su
```

Password:

```
root@seed-desktop:/home/seed# sysctl -w kernel.randomize_va_space=0
```

```
kernel.randomize_va_space = 0 //turn off address randomization
```

```
root@seed-desktop:/home/seed# gcc -o malin malinput.c
```

```
root@seed-desktop:/home/seed# ./malin // file input is generated
```

```
root@seed-desktop:/home/seed# gcc -g -o bof bufOverflow.c
```

```
root@seed-desktop:/home/seed# ./bof < input
Enter the data
Should not reach here
```

```
//bof takes the malicious dump in input file
// as an input
```

The data in 'str' is the same as the text received from the input file . The length of the string 'str' is greater than the size of the buffer 'buf'. The data in input copied into buf causes a buffer overflow.

The return address of copyData in main, stored on the stack gets replaced by the address of the method message() [This is demonstrated with address values below] . As a result after the copy instead of returning back to main the copyData() method invokes the message() function.

The message() is indirectly called from the buffer overflow in copyData(). If there is sensitive information in message() this specific strncpy call can be very dangerous.

```
root@seed-desktop:/home/seed# gdb bof
```

```
GNU gdb 6.8-debian
```

```
Copyright (C) 2008 Free Software Foundation, Inc.
```

```
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
```

```
This is free software: you are free to change and redistribute it.
```

```
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
```

```
This GDB was configured as "i486-linux-gnu"...
```

```
(gdb) list copyData
```

```
7      puts("Should not reach here");
8      exit(0);
9  }
```

```
10
```

```
11      void copyData(char *str)                // I initialized buf to BBB to observe change of value
12      {                                        // of x buf from 0x00424242 to 0x41414141 right after
13          char buf[4] = "BBB";                //strncpy()
14          strncpy(buf, str, strlen(str));
15      }
16
```

```
(gdb) b 13                                     // breakpoint set at first line of method copyData()
```

```
Breakpoint 1 at 0x8048508: file bufOverflow.c, line 13.
```

```
(gdb) run < input
```

```
Starting program: /home/seed/bof < input
```

```
Enter the data
```

```
Breakpoint 1, copyData (str=0xbffff514 "AAAAAAAAA\204\004\b")
at bufOverflow.c:13
```

```
13      char buf[4] = "BBB";
```

```
(gdb) X/i $pc
```

```
0x8048508 <copyData+6>:      mov     0x8048676,%eax
```

```
(gdb) si
```

```
0x0804850d 13      char buf[4] = "BBB";
```

```
(gdb) si
```

```
14      strncpy(buf, str, strlen(str));
```

```
(gdb) si
```

```
0x08048513 14      strncpy(buf, str, strlen(str));
```

```
(gdb) si
```

```
0x08048516 14      strncpy(buf, str, strlen(str));
```

```

(gdb) X/i $pc
0x8048516 <copyData+20>:  call 0x80483f0 <strlen@plt>
(gdb) nexti // nexti called to avoid stepping into strlen
0x0804851b 14  strcpy(buf, str, strlen(str));
(gdb) X/i $pc
0x804851b <copyData+25>:  mov %eax,0x8(%esp)
(gdb) si
0x0804851f 14  strcpy(buf, str, strlen(str));
(gdb) si
0x08048522 14  strcpy(buf, str, strlen(str));
(gdb) si
0x08048526 14  strcpy(buf, str, strlen(str)); // calls before strcpy() will execute
(gdb) si
0x08048529 14  strcpy(buf, str, strlen(str));
(gdb) si
0x0804852c 14  strcpy(buf, str, strlen(str));
(gdb) X/i $pc // The point of debug right before the call to strcpy()
0x804852c <copyData+42>:  call 0x80483c0 <strcpy@plt>
(gdb) disas copyData

```

// Observing the assembly code of copyData function

Dump of assembler code for function copyData:

```

0x08048502 <copyData+0>:  push %ebp
0x08048503 <copyData+1>:  mov %esp,%ebp
0x08048505 <copyData+3>:  sub $0x28,%esp // prolog ends here
0x08048508 <copyData+6>:  mov 0x8048676,%eax
0x0804850d <copyData+11>: mov %eax,-0x4(%ebp)
0x08048510 <copyData+14>: mov 0x8(%ebp),%eax
0x08048513 <copyData+17>: mov %eax,(%esp)
0x08048516 <copyData+20>: call 0x80483f0 <strlen@plt>
0x0804851b <copyData+25>: mov %eax,0x8(%esp)
0x0804851f <copyData+29>: mov 0x8(%ebp),%eax
0x08048522 <copyData+32>: mov %eax,0x4(%esp)
0x08048526 <copyData+36>: lea -0x4(%ebp),%eax
0x08048529 <copyData+39>: mov %eax,(%esp)
0x0804852c <copyData+42>: call 0x80483c0 <strcpy@plt> // Currently $pc is pointing here
0x08048531 <copyData+47>: leave // i.e. just before strcpy executes
0x08048532 <copyData+48>: ret
End of assembler dump. // leave ret are the epilog

```

Stack just before the strcpy() executes :

❖ buf and \$ebp-4 value

```

(gdb) x/s buf
0xbffff4f4:  "BBB"
(gdb) x/x $ebp-4
0xbffff4f4:  0x00424242
(gdb) x/s $ebp-4 // $ebp-4 stores initialized value of 'buf'.
0xbffff4f4:  "BBB"

```

❖ \$ebp

(gdb) x/x \$ebp

0xbffff4f8: 0xbffff528

// address of new ebp is 0xbffff4f8 storing old ebp address 0xbffff528

❖ \$ebp+4

(gdb) x/x \$ebp+4

0xbffff4fc: 0x0804857f

// return address of copyData in main stored at \$ebp+4

(gdb) x/s str

0xbffff514: "AAAAAAA\204\004\b"

(gdb) x/x \$esp

// the input file data is in 'str'

0xbffff4d0: 0xbffff4f4

(gdb) x/x \$esp+4

0xbffff4d4: 0xbffff514

// Observing the assembly code of main

(gdb) disas main

Dump of assembler code for function main:

```
0x08048533 <main+0>:    lea  0x4(%esp),%ecx
0x08048537 <main+4>:    and  $0xffffffff,%esp
0x0804853a <main+7>:    pushl -0x4(%ecx)
0x0804853d <main+10>:   push  %ebp
0x0804853e <main+11>:   mov  %esp,%ebp
0x08048540 <main+13>:   push  %ecx
0x08048541 <main+14>:   sub  $0x24,%esp
0x08048544 <main+17>:   mov  0x4(%ecx),%eax
0x08048547 <main+20>:   mov  %eax,-0x18(%ebp)
0x0804854a <main+23>:   mov  %gs:0x14,%eax
0x08048550 <main+29>:   mov  %eax,-0x8(%ebp)
0x08048553 <main+32>:   xor  %eax,%eax
0x08048555 <main+34>:   movl  $0x804867a,(%esp)
0x0804855c <main+41>:   call 0x8048410 <puts@plt>
0x08048561 <main+46>:   lea  -0x14(%ebp),%eax
0x08048564 <main+49>:   mov  %eax,0x4(%esp)
0x08048568 <main+53>:   movl  $0x8048689,(%esp)
0x0804856f <main+60>:   call 0x80483e0 <scanf@plt>
0x08048574 <main+65>:   lea  -0x14(%ebp),%eax
0x08048577 <main+68>:   mov  %eax,(%esp)
0x0804857a <main+71>:   call 0x8048502 <copyData>
0x0804857f <main+76>:   mov  $0x0,%eax           // return address of copyData in main
0x08048584 <main+81>:   mov  -0x8(%ebp),%edx
0x08048587 <main+84>:   xor  %gs:0x14,%edx
0x0804858e <main+91>:   je   0x8048595 <main+98>
0x08048590 <main+93>:   call 0x8048400 <__stack_chk_fail@plt>
0x08048595 <main+98>:   add  $0x24,%esp
0x08048598 <main+101>:  pop  %ecx
0x08048599 <main+102>:  pop  %ebp
0x0804859a <main+103>:  lea  -0x4(%ecx),%esp
0x0804859d <main+106>:  ret
End of assembler dump.
```

```
❖ (gdb) x/x $ebp+4
0xbffff4fc:    0x0804857f
(gdb) x/x $ebp+8
0xbffff500:    0xbffff514
```

```
// $ebp+4 has same value as the address highlighted in main
// return address of copyData() function is stored at $ebp+4
```

Stack before strcpy()

[esp] 0xbffff4d0	0xbffff4f4	
[esp + 4] 0xbffff4d4	0xbffff514	
[ebp - 4] 0xbffff4f4	0x00424242	[initial value of buf]
[ebp] 0xbffff4f8	0xbffff528	[old value of ebp stored here]
[ebp + 4] 0xbffff4fc	0x0804857f	[return value of copyData in main]
[ebp + 8] 0xbffff500	0xbffff514	
[str] 0xbffff514	"AAAAAAAA[0x20]004[0x00]b"	[value of input string]

```
(gdb) n
```

15 }

```
(gdb) x/i $pc
```

```
// strncpy() function executed
```

```
0x8048531 <copyData+47>:  leave
```

Stack just after the strncpy() executes :

❖ (gdb) x/s buf

0xbfff4f4: "AAAAAAA◆\204\004\b\024◆◆◆\024◆◆◆\030◆◆◆\234\203\004\b◆◆◆AAAAAAA◆\204\004\b"

```
(gdb) x/x $ebp-4
```

```
0xbffff4f4:    0x41414141
```

```
// $ebp-4 stores value of 'buf' after overflow.
```

(gdb) **disas message**

Dump of assembler code for function message:

```
0x080484e4 <message+0>:    push    %ebp                // prolog of the message function
```

```
0x080484e5 <message+1>:    mov    %esp,%ebp
```

```
0x080484e7 <message+3>:    sub    $0x8,%esp
```

```
0x080484ea <message+6>:    movl    $0x8048660,(%esp)    // address of the message function
```

```
0x080484f1 <message+13>:    call 0x8048410 <puts@plt>
```

```
0x080484f6 <message+18>:    movl    $0x0,(%esp)
```

```
0x080484fd <message+25>:    call    0x8048420 <exit@plt>
```

End of assembler dump.

(gdb) x/x \$ebp+4

0xbffff4fc: 0x080484ea

//after strcpy call, \$ebp+4 now stores value of message() address

(gdb) x/x \$ebp

0xbffff4f8: 0x41414141

// \$ebp and \$ebp-4 contain the 'str' contents.

(gdb) x/x \$ebp-8

0xbffff4f0: 0x00000000

(gdb) x/x \$esp

0xbffff4d0: 0xbffff4f4

(gdb) x/x \$esp+4

0xbffff4d4: 0xbffff514

Stack just after strcpy function executes

[esp] 0xbffff4d0	0xbffff4f4
[esp + 4] 0xbffff4d4	0xbffff514
[ebp - 4] 0xbffff4f4	0x41414141
[ebp] 0xbffff4f8	0x41414141
[ebp + 4] 0xbffff4fc	0x080484ea

(c) Rewrite the program using any one of the mitigation strategies that can be used to prevent buffer overflows.

Answer)

❖ **Strncpy()** can be used within copyData as a mitigation strategy, within copyData :
This function is similar to strncpy(), but it copies at most *size-1* bytes to *dest*, always adds a terminating null byte.
The caller must handle the possibility of data loss if the size of *dest* is too small for the source.
The return value of the function is the length of *src*, which allows truncation to be easily detected.

If the return value is greater than or equal to *size* , truncation occurred. If loss of data matters, the caller must either check the arguments before the call, or test the function return value.

The code of bufOverflow can be modified as below :

```
# include <stdio.h>
# include <string.h>
# include <stdlib.h>

#include <sys/types.h>

/*
 * Copy src to string dst of size siz. At most siz-1 characters * will be copied. Always NUL terminates (unless siz
 == 0). * Returns strlen(src); if retval >= siz, truncation occurred. */

size_t strncpy(dst, src, siz)
    char *dst;
    const char *src;
    size_t siz;
{
    register char *d = dst;
    register const char *s = src;
    register size_t n = siz;

    /* Copy as many bytes as will fit */
    if (n != 0 && --n != 0) {
        do {
            if ((*d++ = *s++) == 0)
                break;
        } while (--n != 0);
    }
    /* Not enough room in dst, add NUL and traverse rest of src */
    if (n == 0) {
        if (siz != 0)
            *d = '\0';          /* NUL-terminate dst */
        while (*s++)
            ;
    }
    return(s - src - 1);        /* count does not include NUL */
}
```



```

void message()
{
puts("Should not reach here");
exit(0);
}

void copyData(char *str)
{
char buf[4] = "BBB";
size_t length = strlen(str);
if (length >= 4)
    { printf("\nTruncation occurred: Possible data loss");
      printf("\n buf : %s \n", buf); }
}

int main(int argc , char * argv[])
{
char data[12];
printf("Enter the data\n");
scanf("%s", data);
copyData(data);
return 0;
}

```

```
root@seed-desktop:/home/seed# gcc -g -o bof bufOverflow.c
```

```
root@seed-desktop:/home/seed# ./bof < input
```

Enter the data

Truncation occurred: Possible data loss

buf : AAA

This mitigation has checked if the copy done into buf in copyData is valid and preventing buffer overflow.

A message if the length of source is greater than destination is given to alert user that the destination will have loss of data. Although buf has a size of 4 just three A's get copied as the last byte contains the null terminator.

This is a safe method to use to avoid buffer overflow because size of destination is passed as an input. If handled correctly and right size of the destination buffer is known this is a safe method and buffer overflow mitigation strategy.
